

# Assignment 6

Team members: Subhrajit (23210106), Pratyush (23210075)

---

Github Repo URL: [Github Link](#)

To run the code please follow the instructions given in [readme.md](#)

---

## Part I:

### Question a) Implementation:

#### Topology setup:

We've created three routers: ra, rb, and rc; six hosts: h1, h2, h3, h4, h5, and h6; four switches: s1, s2, s3, and s4.

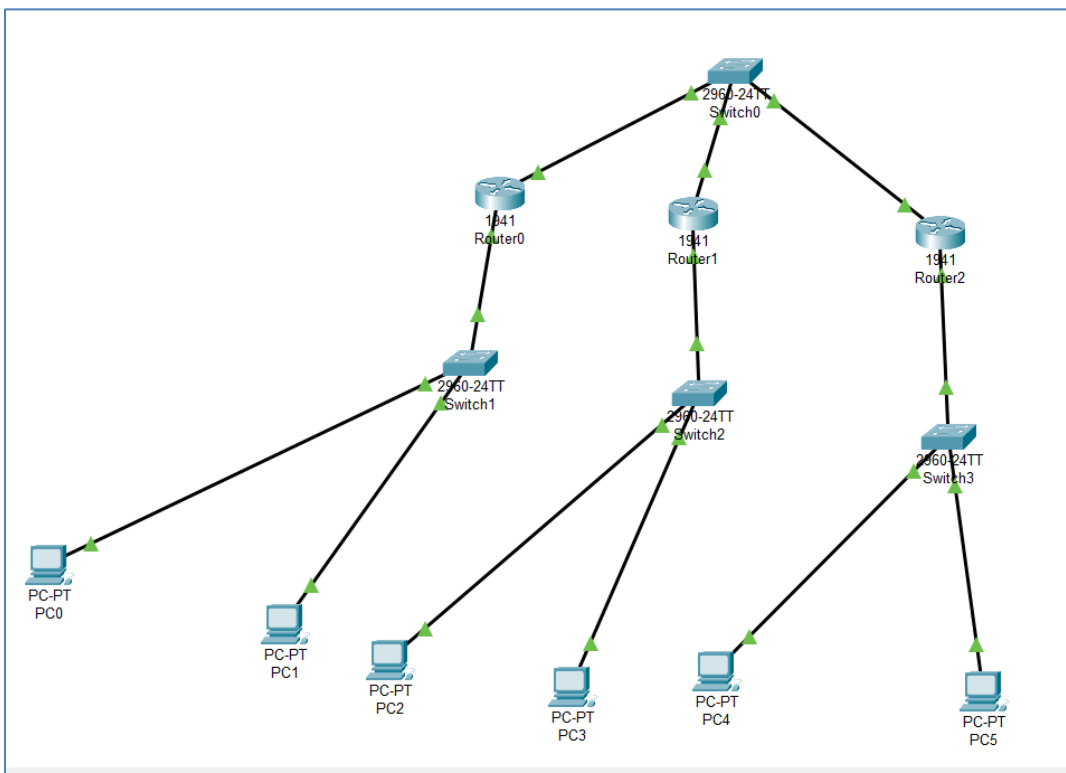
Hosts h1 and h2 are connected to ra via switch s1.

Hosts h3 and h4 are connected to rb via switch s2.

Hosts h5 and h6 are connected to rc via switch s3.

Routers ra, rb, and rc are connected to each other via switch s4.

Topology structure can be visualized like the following:



## IP Configurations:

- 192.168.1.0/24 (ra-h1-h2, IP: 192.168.1.{1, 100, 101})
- 192.168.2.0/24 (rb-h3-h4, IP: 192.168.2.{1, 100, 101})
- 192.168.3.0/24 (rc-h5-h6, IP: 192.168.3.{1, 100, 101})
- 10.0.0.0/8 (ra-rb-rc, IP: 10.0.0.{1, 2, 3})

**Code:** Please see the code for part I implementation here: [topology\\_part1.py](#)

(linked files to follow: [readme.md](#) ; [CLI\\_scripts.md](#) )

(Note: the [topology\\_part1.py](#) file contains only the code to create the network topology. We have to run the CLI scripts written in the [CLI\\_scripts.md](#) to execute according to question a, b, and c.)

## Proof of working of the network topology:

Successful establishment of routers, switches, and hosts:

```
mininet@mininet-vm: ~/Downloads
mininet@mininet-vm:~$ cd Downloads/
mininet@mininet-vm:~/Downloads$ sudo python3 topology_part1.py
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4 h5 h6 ra rb rc
*** Adding switches:
s1 s2 s3 s4
*** Adding links:
(h1, s1) (h2, s1) (h3, s2) (h4, s2) (h5, s3) (h6, s3) (s1, ra) (s2, rb) (s3, rc)
(s4, ra) (s4, rb) (s4, rc)
*** Configuring hosts
h1 h2 h3 h4 h5 h6 ra rb rc
*** Starting controller
c0
*** Starting 4 switches
s1 s2 s3 s4 ...
*** Waiting for switches to connect
s1 s2 s3 s4
*** Routing Table on Routers:
Router ra:
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
0.0.0.0          0.0.0.0         0.0.0.0         U        0      0        0 ra-eth0
10.0.0.0         0.0.0.0         0.0.0.0         U        0      0        0 ra-eth1
192.168.1.0      0.0.0.0         0.0.0.0         U        0      0        0 ra-eth0
192.168.2.0      0.0.0.0         0.0.0.0         U        0      0        0 ra-eth1
192.168.3.0      0.0.0.0         0.0.0.0         U        0      0        0 ra-eth2
```

Ping between all the hosts and routers to test if the network works properly:

```
mininet@mininet-vm: ~/Downloads
192.168.2.0      0.0.0.0         255.255.255.0   U        0      0        0 rb-eth0
192.168.3.0      0.0.0.0         255.255.255.0   UG       0      0        0 rb-eth1
Router rc:
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
0.0.0.0          0.0.0.0         0.0.0.0         U        0      0        0 rc-eth1
192.168.1.0      10.0.0.1        255.255.255.0   UG       0      0        0 rc-eth1
192.168.2.0      10.0.0.2        255.255.255.0   UG       0      0        0 rc-eth1
192.168.3.0      0.0.0.0         255.255.255.0   U        0      0        0 rc-eth0
Run the commands as per CLI_scripts.md :
*** Starting CLI:
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 h3 h4 h5 h6 ra rb rc
h2 -> h1 h3 h4 h5 h6 ra rb rc
h3 -> h1 h2 h4 h5 h6 ra rb rc
h4 -> h1 h2 h3 h5 h6 ra rb rc
h5 -> h1 h2 h3 h4 h6 ra rb rc
h6 -> h1 h2 h3 h4 h5 ra rb rc
ra -> h1 h2 h3 h4 h5 h6 rb rc
rb -> h1 h2 h3 h4 h5 h6 ra rc
rc -> h1 h2 h3 h4 h5 h6 ra rb
*** Results: 0% dropped (72/72 received)
mininet>
```

Dumping:

```
mininet@mininet-vm: ~/Downloads
10.0.0.0      0.0.0.0      255.0.0.0    U      0      0      0 rc-eth1
192.168.1.0   10.0.0.1     255.255.255.0 UG     0      0      0 rc-eth1
192.168.2.0   10.0.0.2     255.255.255.0 UG     0      0      0 rc-eth1
192.168.3.0   0.0.0.0      255.255.255.0 U      0      0      0 rc-eth0
Run the commands as per CLI_scripts.md :
*** Starting CLI:
mininet> pingall
*** Ping: testing ping reachability
h1 -> h2 h3 h4 h5 h6 ra rb rc
h2 -> h1 h3 h4 h5 h6 ra rb rc
h3 -> h1 h2 h4 h5 h6 ra rb rc
h4 -> h1 h2 h3 h5 h6 ra rb rc
h5 -> h1 h2 h3 h4 h6 ra rb rc
h6 -> h1 h2 h3 h4 h5 ra rb rc
ra -> h1 h2 h3 h4 h5 h6 rb rc
rb -> h1 h2 h3 h4 h5 h6 ra rc
rc -> h1 h2 h3 h4 h5 h6 ra rb
*** Results: 0% dropped (72/72 received)
mininet> dump
<Host h1: h1-eth0:192.168.1.100 pid=7962>
<Host h2: h2-eth0:192.168.1.101 pid=7964>
<Host h3: h3-eth0:192.168.2.100 pid=7966>
<Host h4: h4-eth0:192.168.2.101 pid=7968>
<Host h5: h5-eth0:192.168.3.100 pid=7970>
<Host h6: h6-eth0:192.168.3.101 pid=7972>
<LinuxRouter ra: ra-eth0:192.168.1.1,ra-eth1:10.0.0.1 pid=7976>
<LinuxRouter rb: rb-eth0:192.168.2.1,rb-eth1:10.0.0.2 pid=7978>
<LinuxRouter rc: rc-eth0:192.168.3.1,rc-eth1:10.0.0.3 pid=7980>
<OVSSwitch s1: lo:127.0.0.1,s1-eth1:None,s1-eth2:None,s1-eth3:None pid=7985>
<OVSSwitch s2: lo:127.0.0.1,s2-eth1:None,s2-eth2:None,s2-eth3:None pid=7988>
<OVSSwitch s3: lo:127.0.0.1,s3-eth1:None,s3-eth2:None,s3-eth3:None pid=7991>
<OVSSwitch s4: lo:127.0.0.1,s4-eth1:None,s4-eth2:None,s4-eth3:None pid=7994>
<Controller c0: 127.0.0.1:6653 pid=7955>
mininet> 
```

## Question b) Observations:

Pinging 4 packets from h1 to h3 and h1 to h6.

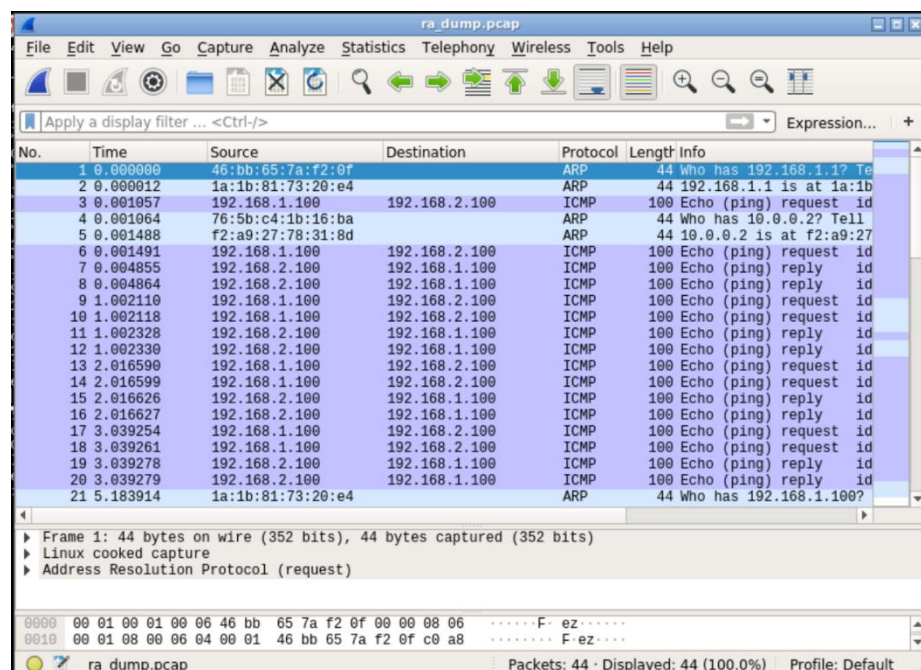
```
mininet@mininet-vm: ~/Downloads
rb -> h1 h2 h3 h4 h5 h6 ra rc
rc -> h1 h2 h3 h4 h5 h6 ra rb
*** Results: 0% dropped (72/72 received)
mininet> h1 ping -c 4 h3
PING 192.168.2.100 (192.168.2.100) 56(84) bytes of data.
64 bytes from 192.168.2.100: icmp_seq=1 ttl=62 time=2.51 ms
64 bytes from 192.168.2.100: icmp_seq=2 ttl=62 time=0.619 ms
64 bytes from 192.168.2.100: icmp_seq=3 ttl=62 time=0.181 ms
64 bytes from 192.168.2.100: icmp_seq=4 ttl=62 time=0.050 ms

--- 192.168.2.100 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3034ms
rtt min/avg/max/mdev = 0.050/0.841/2.517/0.990 ms
mininet> h1 ping -c 4 h6
PING 192.168.3.101 (192.168.3.101) 56(84) bytes of data.
64 bytes from 192.168.3.101: icmp_seq=1 ttl=62 time=3.11 ms
64 bytes from 192.168.3.101: icmp_seq=2 ttl=62 time=0.862 ms
64 bytes from 192.168.3.101: icmp_seq=3 ttl=62 time=0.194 ms
64 bytes from 192.168.3.101: icmp_seq=4 ttl=62 time=0.053 ms

--- 192.168.3.101 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3051ms
rtt min/avg/max/mdev = 0.053/1.056/3.117/1.228 ms
mininet> exit
```

Opened the tcpdump through wireshark at routers ra, rb and rc to check the route setup for both pings.

ra\_dump.pcap:



As we've pinged h1 to h3 first, it starts by ARP for asking 'who has 192.168.1.1?' and gets the ARP reply of MAC address. It then starts looking '10.0.0.2' through ARP. After that it starts the packet transmission from 192.168.1.100 (h1) to 192.168.2.100 (h3) while getting the replies from 192.168.2.100 to 192.168.1.100 via router ra.

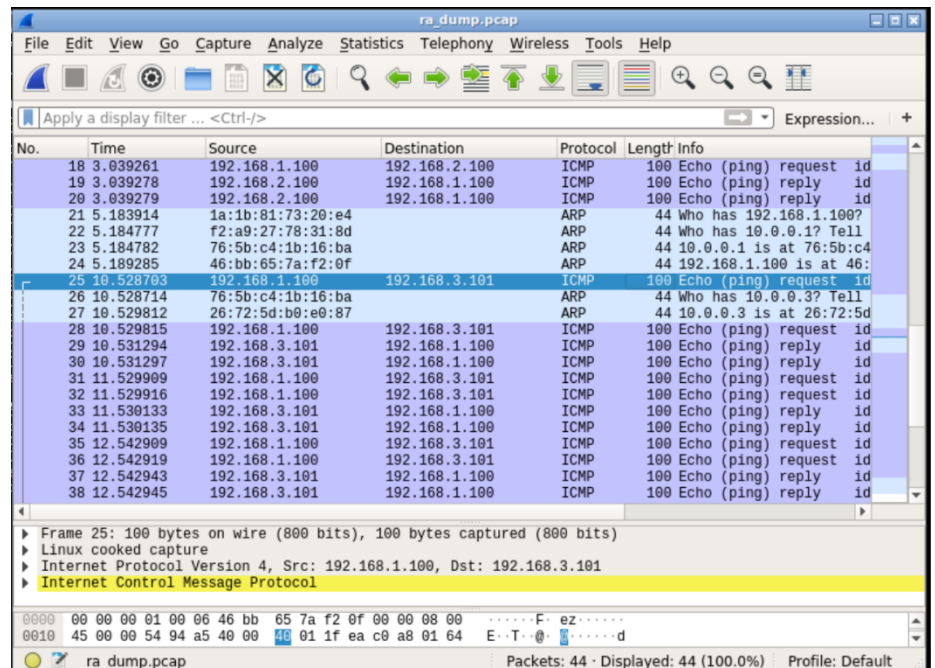
Afterwards for pinging from h1 to h6, we see that in a similar fashion at router ra:

That is,

ICMP packets transmission

Between 192.168.1.100 and

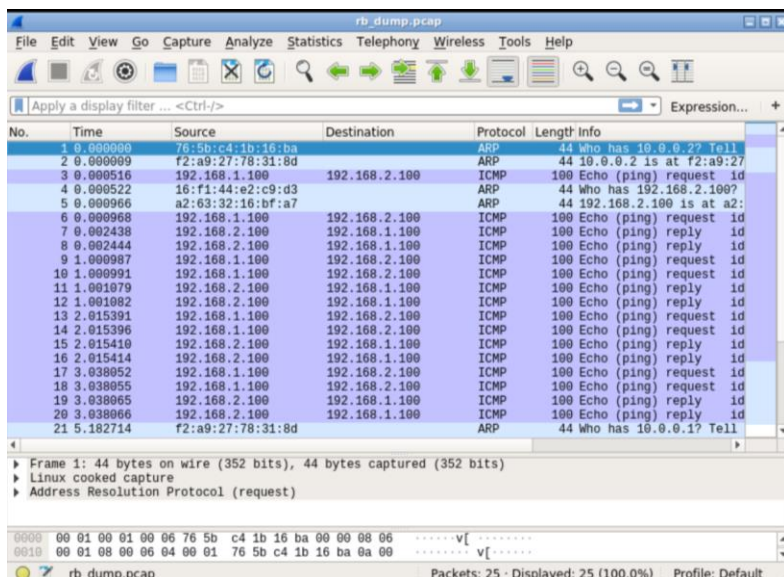
192.168.3.101



The screenshot shows a Wireshark capture of network traffic at router 'ra'. The main display area shows a list of 38 captured packets. The selected packet (No. 25) is an ICMP Echo (ping) request from 192.168.1.100 to 192.168.3.101. The packet details pane shows the Internet Control Message Protocol (ICMP) section. The packet bytes pane shows the raw data in hexadecimal and ASCII.

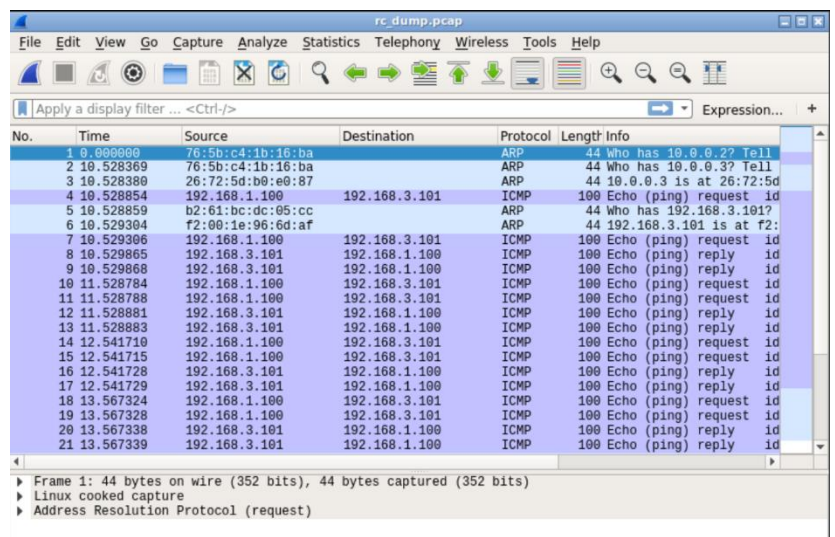
No.	Time	Source	Destination	Protocol	Length	Info
18	3.039261	192.168.1.100	192.168.2.100	ICMP	100	Echo (ping) request id
19	3.039278	192.168.2.100	192.168.1.100	ICMP	100	Echo (ping) reply id
20	3.039279	192.168.2.100	192.168.1.100	ICMP	100	Echo (ping) reply id
21	5.183914	1a:1b:81:73:20:e4		ARP	44	Who has 192.168.1.100?
22	5.184777	f2:a9:27:78:31:8d		ARP	44	Who has 192.168.1.100? Tell
23	5.184782	76:5b:c4:1b:16:ba		ARP	44	192.168.1.100 is at 76:5b:c4
24	5.189285	46:bb:65:7a:f2:0f		ARP	44	192.168.1.100 is at 46:
25	10.528703	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
26	10.528714	76:5b:c4:1b:16:ba		ARP	44	Who has 192.168.3.101? Tell
27	10.529812	26:72:5d:b0:e0:87		ARP	44	192.168.3.101 is at 26:72:5d
28	10.529815	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
29	10.531294	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id
30	10.531297	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id
31	11.529909	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
32	11.529916	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
33	11.530133	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id
34	11.530135	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id
35	12.542909	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
36	12.542919	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
37	12.542943	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id
38	12.542945	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id

You may look at the tcpdump at the routers rb and rc here:



The screenshot shows a Wireshark capture of network traffic at router 'rb'. The main display area shows a list of 21 captured packets. The selected packet (No. 1) is an ARP request from 76:5b:c4:1b:16:ba to 192.168.2.100. The packet details pane shows the Address Resolution Protocol (request) section. The packet bytes pane shows the raw data in hexadecimal and ASCII.

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	76:5b:c4:1b:16:ba		ARP	44	Who has 192.168.2.100? Tell
2	0.000000	f2:a9:27:78:31:8d		ARP	44	Who has 192.168.2.100? Tell
3	0.000516	192.168.1.100	192.168.2.100	ICMP	100	Echo (ping) request id
4	0.000522	16:f1:44:e2:c9:d3		ARP	44	Who has 192.168.2.100?
5	0.000966	a2:63:32:10:bf:a7		ARP	44	192.168.2.100 is at a2:
6	0.000968	192.168.1.100	192.168.2.100	ICMP	100	Echo (ping) request id
7	0.002438	192.168.2.100	192.168.1.100	ICMP	100	Echo (ping) reply id
8	0.002444	192.168.2.100	192.168.1.100	ICMP	100	Echo (ping) reply id
9	1.000987	192.168.1.100	192.168.2.100	ICMP	100	Echo (ping) request id
10	1.000991	192.168.1.100	192.168.2.100	ICMP	100	Echo (ping) request id
11	1.001079	192.168.2.100	192.168.1.100	ICMP	100	Echo (ping) reply id
12	1.001082	192.168.2.100	192.168.1.100	ICMP	100	Echo (ping) reply id
13	2.015391	192.168.1.100	192.168.2.100	ICMP	100	Echo (ping) request id
14	2.015396	192.168.1.100	192.168.2.100	ICMP	100	Echo (ping) request id
15	2.015410	192.168.2.100	192.168.1.100	ICMP	100	Echo (ping) reply id
16	2.015414	192.168.2.100	192.168.1.100	ICMP	100	Echo (ping) reply id
17	3.038052	192.168.1.100	192.168.2.100	ICMP	100	Echo (ping) request id
18	3.038055	192.168.1.100	192.168.2.100	ICMP	100	Echo (ping) request id
19	3.038065	192.168.2.100	192.168.1.100	ICMP	100	Echo (ping) reply id
20	3.038066	192.168.2.100	192.168.1.100	ICMP	100	Echo (ping) reply id
21	5.182714	f2:a9:27:78:31:8d		ARP	44	Who has 192.168.2.100? Tell



The screenshot shows a Wireshark capture of network traffic at router 'rc'. The main display area shows a list of 21 captured packets. The selected packet (No. 1) is an ARP request from 76:5b:c4:1b:16:ba to 192.168.3.101. The packet details pane shows the Address Resolution Protocol (request) section. The packet bytes pane shows the raw data in hexadecimal and ASCII.

No.	Time	Source	Destination	Protocol	Length	Info
1	0.000000	76:5b:c4:1b:16:ba		ARP	44	Who has 192.168.3.101? Tell
2	10.528369	76:5b:c4:1b:16:ba		ARP	44	Who has 192.168.3.101? Tell
3	10.528390	26:72:5d:b0:e0:87		ARP	44	192.168.3.101 is at 26:72:5d
4	10.528854	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
5	10.528859	b2:61:bc:dc:05:cc		ARP	44	Who has 192.168.3.101?
6	10.529304	f2:00:1e:96:6d:af		ARP	44	192.168.3.101 is at f2:
7	10.529306	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
8	10.529865	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id
9	10.529868	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id
10	11.528784	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
11	11.528788	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
12	11.528881	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id
13	11.528883	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id
14	12.541710	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
15	12.541715	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
16	12.541728	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id
17	12.541729	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id
18	13.567324	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
19	13.567328	192.168.1.100	192.168.3.101	ICMP	100	Echo (ping) request id
20	13.567338	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id
21	13.567339	192.168.3.101	192.168.1.100	ICMP	100	Echo (ping) reply id



### Question c) Comparing latency differences between default and changed route.

With the default route:

```
Run the commands as per CLI_scripts.md :
*** Starting CLI:
mininet> h1 traceroute h6
traceroute to 192.168.3.101 (192.168.3.101), 30 hops max, 60 byte packets
 1 192.168.1.1 (192.168.1.1) 324.985 ms 342.809 ms 343.090 ms
 2 10.0.0.3 (10.0.0.3) 347.403 ms 348.209 ms 348.368 ms
 3 192.168.3.101 (192.168.3.101) 352.957 ms 353.240 ms 353.355 ms
mininet>
```

As we can see the route for h1 (192.168.1.101) to h6 (192.168.3.101) goes directly via 10.0.0.3 (rc). That is, it takes the path from h1→ra→rc→h6

After we change the route of h1 to h6 via rb:

```
mininet> ra ip route
10.0.0.0/8 dev ra-eth1 proto kernel scope link src 10.0.0.1
192.168.1.0/24 dev ra-eth0 proto kernel scope link src 192.168.1.1
192.168.2.0/24 via 10.0.0.2 dev ra-eth1
192.168.3.0/24 via 10.0.0.3 dev ra-eth1
mininet> ra ip route del 192.168.3.0/24
mininet> ra ip route add 192.168.3.0/24 via 10.0.0.2 dev ra-eth1
mininet> ra ip route
10.0.0.0/8 dev ra-eth1 proto kernel scope link src 10.0.0.1
192.168.1.0/24 dev ra-eth0 proto kernel scope link src 192.168.1.1
192.168.2.0/24 via 10.0.0.2 dev ra-eth1
192.168.3.0/24 via 10.0.0.2 dev ra-eth1
mininet> h1 traceroute h6
traceroute to 192.168.3.101 (192.168.3.101), 30 hops max, 60 byte packets
 1 192.168.1.1 (192.168.1.1) 111.385 ms 123.947 ms 123.972 ms
 2 10.0.0.2 (10.0.0.2) 126.347 ms 126.355 ms 126.364 ms
 3 10.0.0.3 (10.0.0.3) 129.318 ms 129.324 ms 129.346 ms
 4 192.168.3.101 (192.168.3.101) 130.151 ms 130.158 ms 130.160 ms
mininet>
```

As you can see router ra has the route entry to rc connected directly via 10.0.0.3.

We modified this entry so that ra has the route entry to rc via 10.0.0.2 that is through rb.

The traceroute packet also shows that path from h1 to h6 is now h1→ra→rb→rc→h6.

```

Run the commands as per CLI_scripts.md :
*** Starting CLI:
mininet> h1 ping -c 4 h6
PING 192.168.3.101 (192.168.3.101) 56(84) bytes of data.
64 bytes from 192.168.3.101: icmp_seq=1 ttl=62 time=3.35 ms
64 bytes from 192.168.3.101: icmp_seq=2 ttl=62 time=0.399 ms
64 bytes from 192.168.3.101: icmp_seq=3 ttl=62 time=0.075 ms
64 bytes from 192.168.3.101: icmp_seq=4 ttl=62 time=0.057 ms

--- 192.168.3.101 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3049ms
rtt min/avg/max/mdev = 0.057/0.972/3.357/1.383 ms
mininet> iperf h1 h6
*** Iperf: testing TCP bandwidth between h1 and h6
*** Results: ['10.3 Gbits/sec', '10.3 Gbits/sec']
mininet> ra ip route del 192.168.3.0/24
mininet> ra ip route add 192.168.3.0/24 via 10.0.0.2 dev ra-eth1
mininet> h1 ping -c 4 h6
PING 192.168.3.101 (192.168.3.101) 56(84) bytes of data.
64 bytes from 192.168.3.101: icmp_seq=1 ttl=62 time=5.23 ms
64 bytes from 192.168.3.101: icmp_seq=2 ttl=62 time=0.236 ms
64 bytes from 192.168.3.101: icmp_seq=3 ttl=62 time=0.061 ms
64 bytes from 192.168.3.101: icmp_seq=4 ttl=62 time=0.094 ms

--- 192.168.3.101 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3034ms
rtt min/avg/max/mdev = 0.061/1.405/5.230/2.209 ms
mininet> iperf h1 h6
*** Iperf: testing TCP bandwidth between h1 and h6
*** Results: ['8.83 Gbits/sec', '8.85 Gbits/sec']
mininet>

```

So, if we compare the default route and updated route:

Route	Ping (rtt min/avg/max/mdev)	Iperf (Bandwidth)
Default (h → ra → rc → h6)	0.057/0.972/3.357/1.383 ms	10.3 Gb/s, 10.3 Gb/s
Modified (h1 → ra → rb → rc → h6)	0.061/1.405/5.230/2.209 ms	8.83 Gb/s, 8.85 Gb/s

#### Question d) Dumping routing tables:

Routing tables in context of question b)

```
mininet@mininet-vm: ~/Downloads
4 packets transmitted, 4 received, 0% packet loss, time 3043ms
rtt min/avg/max/mdev = 0.049/1.603/5.939/2.506 ms
mininet> h1 ping -c 4 h6
PING 192.168.3.101 (192.168.3.101) 56(84) bytes of data:
64 bytes from 192.168.3.101: icmp_seq=1 ttl=62 time=17.9 ms
64 bytes from 192.168.3.101: icmp_seq=2 ttl=62 time=0.488 ms
64 bytes from 192.168.3.101: icmp_seq=3 ttl=62 time=0.129 ms
64 bytes from 192.168.3.101: icmp_seq=4 ttl=62 time=0.062 ms

--- 192.168.3.101 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3045ms
rtt min/avg/max/mdev = 0.062/4.665/17.984/7.691 ms
mininet> ra route
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
10.0.0.0          0.0.0.0         255.0.0.0       U        0      0        0 ra-eth1
192.168.1.0       0.0.0.0         255.255.255.0   U        0      0        0 ra-eth0
192.168.2.0       10.0.0.2        255.255.255.0   UG       0      0        0 ra-eth1
192.168.3.0       10.0.0.3        255.255.255.0   UG       0      0        0 ra-eth1
mininet> rb route
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
10.0.0.0          0.0.0.0         255.0.0.0       U        0      0        0 rb-eth1
192.168.1.0       10.0.0.1        255.255.255.0   UG       0      0        0 rb-eth1
192.168.2.0       0.0.0.0         255.255.255.0   U        0      0        0 rb-eth0
192.168.3.0       10.0.0.3        255.255.255.0   UG       0      0        0 rb-eth1
mininet> rc route
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
10.0.0.0          0.0.0.0         255.0.0.0       U        0      0        0 rc-eth1
192.168.1.0       10.0.0.1        255.255.255.0   UG       0      0        0 rc-eth1
192.168.2.0       10.0.0.2        255.255.255.0   UG       0      0        0 rc-eth1
192.168.3.0       0.0.0.0         255.255.255.0   U        0      0        0 rc-eth0
mininet>
```



## Routing tables in context of question c)

```

mininet@mininet-vm: ~/Downloads
PING 192.168.3.101 (192.168.3.101) 56(84) bytes of data.
64 bytes from 192.168.3.101: icmp_seq=1 ttl=62 time=5.23 ms
64 bytes from 192.168.3.101: icmp_seq=2 ttl=62 time=0.236 ms
64 bytes from 192.168.3.101: icmp_seq=3 ttl=62 time=0.061 ms
64 bytes from 192.168.3.101: icmp_seq=4 ttl=62 time=0.094 ms

--- 192.168.3.101 ping statistics ---
4 packets transmitted, 4 received, 0% packet loss, time 3034ms
rtt min/avg/max/mdev = 0.061/1.405/5.230/2.209 ms
mininet> iperf h1 h6
*** Iperf: testing TCP bandwidth between h1 and h6
*** Results: ['8.83 Gbits/sec', '8.85 Gbits/sec']
mininet> ra route
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
10.0.0.0          0.0.0.0         255.0.0.0       U        0      0        0 ra-eth1
192.168.1.0       0.0.0.0         255.255.255.0   U        0      0        0 ra-eth0
192.168.2.0       10.0.0.2        255.255.255.0   UG       0      0        0 ra-eth1
192.168.3.0       10.0.0.2        255.255.255.0   UG       0      0        0 ra-eth1
mininet> rb route
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
10.0.0.0          0.0.0.0         255.0.0.0       U        0      0        0 rb-eth1
192.168.1.0       10.0.0.1        255.255.255.0   UG       0      0        0 rb-eth1
192.168.2.0       0.0.0.0         255.255.255.0   U        0      0        0 rb-eth0
192.168.3.0       10.0.0.3        255.255.255.0   UG       0      0        0 rb-eth1
mininet> rc route
Kernel IP routing table
Destination      Gateway         Genmask         Flags Metric Ref    Use Iface
10.0.0.0          0.0.0.0         255.0.0.0       U        0      0        0 rc-eth1
192.168.1.0       10.0.0.1        255.255.255.0   UG       0      0        0 rc-eth1
192.168.2.0       10.0.0.2        255.255.255.0   UG       0      0        0 rc-eth1
192.168.3.0       0.0.0.0         255.255.255.0   U        0      0        0 rc-eth0
mininet>

```

## Part II

Question **a)** Implementation of mininet topology and TCP-Server

Created four hosts: h1, h2, h3 and h4; and two switches: s1 and s2.

Hosts h1 and h2 are connected to s1; and h3 and h4 are connected to s2.

s1 and s2 are connected to each other.

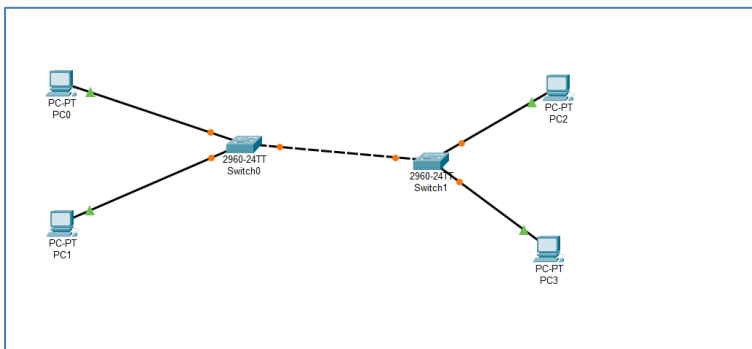
IP of the hosts: h1: 192.168.1.1/24

h2: 192.168.1.2/24

h3: 192.168.1.3/24

h4: 192.168.1.4/24

Visualization:



**Code:** Please see the code for part II implementation here: [topology\\_part2.py](#)

(please follow the instructions given here: [readme.md](#) for part II)

Proof of working of the network:

(executed with -- config b, no loss and congestion control scheme reno)

```
mininet@mininet-vm: ~/Downloads
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2
*** Adding links:
(2,00Mbit 2 delay) (2,00Mbit 2 delay) (h1, s1) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h2, s1) (2,00Mbit 2 delay)
(2,00Mbit 2 delay) (h3, s2) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h4, s2) (8,00Mbit 2 delay 0,00000% loss) (8,00
Mbit 2 delay 0,00000% loss) (s1, s2)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0
*** Starting 2 switches
s1 s2 ... (2,00Mbit 2 delay) (2,00Mbit 2 delay) (8,00Mbit 2 delay 0,00000% loss) (2,00Mbit 2 delay) (2,00Mbit 2 de
lay) (8,00Mbit 2 delay 0,00000% loss)
*** Waiting for switches to connect
s1 s2
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85,3 KByte (default)
-----
[ 3] local 192.168.1.1 port 58654 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0,0-10,3 sec  9,88 MBytes 8,03 Mbits/sec

*** Starting CLI:
mininet> exit
```

Question b)

Running client on h1 and server on h4:

Scheme: reno

```
mininet@mininet-vm: ~/Downloads
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2
*** Adding links:
(2,00Mbit 2 delay) (2,00Mbit 2 delay) (h1, s1) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h2, s1) (2,00Mbit 2 delay)
(2,00Mbit 2 delay) (h3, s2) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h4, s2) (8,00Mbit 2 delay 0,00000% loss) (8,00
Mbit 2 delay 0,00000% loss) (s1, s2)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0
*** Starting 2 switches
s1 s2 ... (2,00Mbit 2 delay) (2,00Mbit 2 delay) (8,00Mbit 2 delay 0,00000% loss) (2,00Mbit 2 delay) (2,00Mbit 2 de
lay) (8,00Mbit 2 delay 0,00000% loss)
*** Waiting for switches to connect
s1 s2
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85,3 KByte (default)
-----
[ 3] local 192.168.1.1 port 58654 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0,0-10,3 sec  9,88 MBytes 8,03 Mbits/sec

*** Starting CLI:
mininet> exit
```

Scheme: vegas

```
mininet@mininet-vm: ~/Downloads
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2
*** Adding links:
(2,00Mbit 2 delay) (2,00Mbit 2 delay) (h1, s1) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h2, s1) (2,00Mbit 2 delay)
(2,00Mbit 2 delay) (h3, s2) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h4, s2) (8,00Mbit 2 delay 0,00000% loss) (8,00
Mbit 2 delay 0,00000% loss) (s1, s2)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0
*** Starting 2 switches
s1 s2 ... (2,00Mbit 2 delay) (2,00Mbit 2 delay) (8,00Mbit 2 delay 0,00000% loss) (2,00Mbit 2 delay) (2,00Mbit 2 de
lay) (8,00Mbit 2 delay 0,00000% loss)
*** Waiting for switches to connect
s1 s2
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85,3 KByte (default)
-----
[ 3] local 192.168.1.1 port 58750 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0,0-10,4 sec  2,62 MBytes 2,13 Mbits/sec

*** Starting CLI:
mininet>
```

## Scheme: cubic

```
mininet@mininet-vm: ~/Downloads
mininet@mininet-vm:~/Downloads$ sudo python3 topology_part2.py --config b --cc cubic
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2
*** Adding links:
(2.00Mbit 2 delay) (2.00Mbit 2 delay) (h1, s1) (2.00Mbit 2 delay) (2.00Mbit 2 delay) (h2, s1) (2.00Mbit 2 delay)
(2.00Mbit 2 delay) (h3, s2) (2.00Mbit 2 delay) (2.00Mbit 2 delay) (h4, s2) (8.00Mbit 2 delay 0.00000% loss) (8.00
Mbit 2 delay 0.00000% loss) (s1, s2)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0
*** Starting 2 switches
s1 s2 ... (2.00Mbit 2 delay) (2.00Mbit 2 delay) (8.00Mbit 2 delay 0.00000% loss) (2.00Mbit 2 delay) (2.00Mbit 2 de
lay) (8.00Mbit 2 delay 0.00000% loss)
*** Waiting for switches to connect
s1 s2

-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 3] local 192.168.1.1 port 58762 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 3] 0.0-10.8 sec  4.38 MBytes  3.39 Mbits/sec

*** Starting CLI:
mininet>
```

## Scheme: bbr

```
mininet@mininet-vm: ~/Downloads
mininet@mininet-vm:~/Downloads$ sudo python3 topology_part2.py --config b --cc bbr
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2
*** Adding links:
(2.00Mbit 2 delay) (2.00Mbit 2 delay) (h1, s1) (2.00Mbit 2 delay) (2.00Mbit 2 delay) (h2, s1) (2.00Mbit 2 delay)
(2.00Mbit 2 delay) (h3, s2) (2.00Mbit 2 delay) (2.00Mbit 2 delay) (h4, s2) (8.00Mbit 2 delay 0.00000% loss) (8.00
Mbit 2 delay 0.00000% loss) (s1, s2)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0
*** Starting 2 switches
s1 s2 ... (2.00Mbit 2 delay) (2.00Mbit 2 delay) (8.00Mbit 2 delay 0.00000% loss) (2.00Mbit 2 delay) (2.00Mbit 2 de
lay) (8.00Mbit 2 delay 0.00000% loss)
*** Waiting for switches to connect
s1 s2

-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 128 KByte (default)
-----
[ 3] local 192.168.1.1 port 58770 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 3] 0.0-10.3 sec  2.50 MBytes  2.04 Mbits/sec

*** Starting CLI:
mininet>
```

## Summary of the throughput: (h1→h4)

Scheme	Throughput	Transferred Data
RENO	8.03 Mbits/sec	9.88 MB
VEGAS	2.13 Mbits/sec	2.62 MB
CUBIC	3.39 Mbits/sec	4.38 MB
BBR	2.04 Mbits/sec	2.50 MB

We're seeing the maximum throughput in reno scheme, whereas the least in bbr scheme.

## Reasoning of the throughput differences:

The observed throughput differences between the Congestion Control Algorithms (CCAs) can be attributed to their distinct mechanisms for dealing with network congestion. Each CCA employs a different approach to adjusting transmission rates in response to network conditions, leading to varying throughput performance.

Here's a breakdown of the CCAs and their impact on throughput:

**Reno:** The default CCA for TCP, Reno is a loss-based algorithm that slows down the transmission rate upon packet loss, assuming network congestion. While simple and easy to implement, Reno's conservative rate reduction can limit throughput in congested networks. But the iperf test methodology might have inadvertently favored Reno leading to the max throughput. Also, it got the highest size of transferred data, that also may be reason of getting higher throughput.

**Vegas:** Vegas utilizes both packet loss and delay information to estimate network congestion. This allows for more accurate rate adjustments compared to Reno, potentially leading to higher throughput in some scenarios. But in this case, the iperf test methodology might have not favoured reno, leading to the lower throughput.

**Cubic:** Cubic employs a combination of loss, delay, and packet arrival times to estimate congestion, enabling quicker rate adjustments compared to Vegas. That may be the reason of it getting higher throughput than vegas.

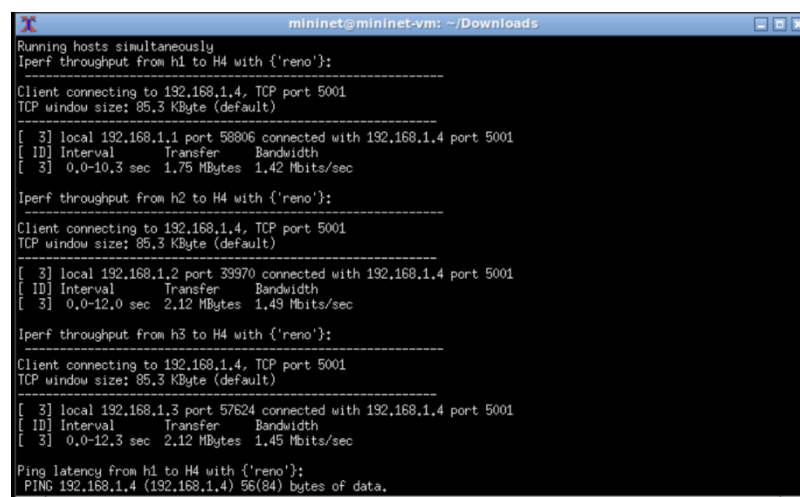
**BBR:** Google's BBR CCA relies on a delay-based approach, using a network model to estimate the optimal sending rate. BBR can achieve high throughput in networks with high latency or packet loss, but it may be more complex to implement than Reno, Vegas, or Cubic. And iperf testing methodology may have negatively impacted the throughput.

### Question c)

To simultaneously run the three clients h1, h2, and h3 on h4 server, we did the following:

We first started iperf server on h4. Then we created three threads for h1, h2, and h3 as clients to be served by h4. Then we simultaneously ran the three threads to be served by h4. And all threads are defined on the run() function during the net active period.

Scheme: reno



```
mininet@mininet-vm: ~/Downloads
Running hosts simultaneously
Iperf throughput from h1 to H4 with ('reno'):
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85,3 KByte (default)
[ 3] local 192.168.1.1 port 58806 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0,0-10,3 sec  1,75 MBytes 1,42 Mbits/sec

Iperf throughput from h2 to H4 with ('reno'):
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85,3 KByte (default)
[ 3] local 192.168.1.2 port 39970 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0,0-12,0 sec  2,12 MBytes 1,49 Mbits/sec

Iperf throughput from h3 to H4 with ('reno'):
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85,3 KByte (default)
[ 3] local 192.168.1.3 port 57624 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0,0-12,3 sec  2,12 MBytes 1,45 Mbits/sec

Ping latency from h1 to H4 with ('reno'):
PING 192.168.1.4 (192.168.1.4) 56(84) bytes of data.
```

## Scheme: vegas

```
mininet@mininet-vm: ~/Downloads
*** Waiting for switches to connect
s1 s2
Running hosts simultaneously
Iperf throughput from h2 to H4 with {'vegas'}:
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 3] local 192.168.1.2 port 40000 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.0 sec  2.12 MBytes 1.78 Mbits/sec
-----
Iperf throughput from h3 to H4 with {'vegas'}:
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 3] local 192.168.1.3 port 57654 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-12.5 sec  2.12 MBytes 1.43 Mbits/sec
-----
Iperf throughput from h1 to H4 with {'vegas'}:
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 3] local 192.168.1.1 port 58836 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-12.7 sec  2.25 MBytes 1.49 Mbits/sec
```

## Scheme: cubic

```
mininet@mininet-vm: ~/Downloads
*** Waiting for switches to connect
s1 s2
Running hosts simultaneously
Iperf throughput from h1 to H4 with {'cubic'}:
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 3] local 192.168.1.1 port 58866 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-12.1 sec  2.12 MBytes 1.48 Mbits/sec
-----
Iperf throughput from h2 to H4 with {'cubic'}:
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 3] local 192.168.1.2 port 40030 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-12.2 sec  2.12 MBytes 1.46 Mbits/sec
-----
Iperf throughput from h3 to H4 with {'cubic'}:
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 3] local 192.168.1.3 port 57684 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-12.9 sec  2.38 MBytes 1.55 Mbits/sec
```

## Scheme: bbr

```
mininet@mininet-vm: ~/Downloads
*** Waiting for switches to connect
s1 s2
Running hosts simultaneously
Iperf throughput from h2 to H4 with {'bbr'}:
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 3] local 192.168.1.2 port 40042 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.2 sec  1.62 MBytes 1.34 Mbits/sec
-----
Iperf throughput from h1 to H4 with {'bbr'}:
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 3] local 192.168.1.1 port 58874 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.3 sec  1.75 MBytes 1.42 Mbits/sec
-----
Iperf throughput from h3 to H4 with {'bbr'}:
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 3] local 192.168.1.3 port 57700 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-12.0 sec  2.12 MBytes 1.49 Mbits/sec
```



## Summary of the schemes:

Scheme	Source	Throughput	File Transferred
reno	h1→h4	1.4 Mb/s	1.7 MB
reno	h2→h4	1.49 Mb/s	2.12 MB
reno	h3→h4	1.45 Mb/s	2.12 MB
vegas	h1→h4	1.49 Mb/s	2.25 MB
vegas	h2→h4	1.78 Mb/s	2.12 MB
vegas	h3→h4	1.43 Mb/s	2.12 MB
cubic	h1→h4	1.48 Mb/s	2.12 MB
cubic	h2→h4	1.46 Mb/s	2.12 MB
cubic	h3→h4	1.55 Mb/s	2.38 MB
bbr	h1→h4	1.42 Mb/s	1.75 MB
bbr	h2→h4	1.34 Mb/s	1.62 MB
bbr	h3→h4	1.49 Mb/s	2.12 MB

We got similar results across all the schemes. Although the CCAs are different but we still got the similar results may be due to the methodology that is implemented internally by iperf.

## Question d)

Running the `--config b` with all the schemes:

Scheme: reno with 1% loss

```

mininet@mininet-vm: ~/Downloads
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2
*** Adding links:
(2.00Mbit 2 delay) (2.00Mbit 2 delay) (h1, s1) (2.00Mbit 2 delay) (2.00Mbit 2 delay) (h2, s1) (2.00Mbit 2 delay)
(2.00Mbit 2 delay) (h3, s2) (2.00Mbit 2 delay) (2.00Mbit 2 delay) (h4, s2) (8.00Mbit 2 delay 1.00000% loss) (8.00
Mbit 2 delay 1.00000% loss) (s1, s2)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0
*** Starting 2 switches
s1 s2 ... (2.00Mbit 2 delay) (2.00Mbit 2 delay) (8.00Mbit 2 delay 1.00000% loss) (2.00Mbit 2 delay) (2.00Mbit 2 de
lay) (8.00Mbit 2 delay 1.00000% loss)
*** Waiting for switches to connect
s1 s2
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)
-----
[ 3] local 192.168.1.1 port 58936 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer      Bandwidth
[ 3] 0.0-11.1 sec  5.25 MBytes  3.96 Mbits/sec
*** Starting CLI:
mininet>

```

## Scheme: reno with 3% loss

```
mininet@mininet-vm: ~/Downloads
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2
*** Adding links:
(2.00Mbit 2 delay) (2.00Mbit 2 delay) (h1, s1) (2.00Mbit 2 delay) (2.00Mbit 2 delay) (h2, s1) (2.00Mbit 2 delay)
(2.00Mbit 2 delay) (h3, s2) (2.00Mbit 2 delay) (2.00Mbit 2 delay) (h4, s2) (8.00Mbit 2 delay 3.00000% loss) (8.00
Mbit 2 delay 3.00000% loss) (s1, s2)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0
*** Starting 2 switches
s1 s2 ... (2.00Mbit 2 delay) (2.00Mbit 2 delay) (8.00Mbit 2 delay 3.00000% loss) (2.00Mbit 2 delay) (2.00Mbit 2 de
lay) (8.00Mbit 2 delay 3.00000% loss)
*** Waiting for switches to connect
s1 s2

Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)

[ 3] local 192.168.1.1 port 58950 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.5 sec   2.50 MBytes  2.00 Mbits/sec

*** Starting CLI:
mininet>
```

## Scheme: vegas with 1% loss

```
mininet@mininet-vm: ~/Downloads
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2
*** Adding links:
(2.00Mbit 2 delay) (2.00Mbit 2 delay) (h1, s1) (2.00Mbit 2 delay) (2.00Mbit 2 delay) (h2, s1) (2.00Mbit 2 delay)
(2.00Mbit 2 delay) (h3, s2) (2.00Mbit 2 delay) (2.00Mbit 2 delay) (h4, s2) (8.00Mbit 2 delay 1.00000% loss) (8.00
Mbit 2 delay 1.00000% loss) (s1, s2)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0
*** Starting 2 switches
s1 s2 ... (2.00Mbit 2 delay) (2.00Mbit 2 delay) (8.00Mbit 2 delay 1.00000% loss) (2.00Mbit 2 delay) (2.00Mbit 2 de
lay) (8.00Mbit 2 delay 1.00000% loss)
*** Waiting for switches to connect
s1 s2

Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)

[ 3] local 192.168.1.1 port 58958 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.3 sec   2.62 MBytes  2.14 Mbits/sec

*** Starting CLI:
mininet>
```

## Scheme: vegas with 3% loss

```
mininet@mininet-vm: ~/Downloads
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2
*** Adding links:
(2.00Mbit 2 delay) (2.00Mbit 2 delay) (h1, s1) (2.00Mbit 2 delay) (2.00Mbit 2 delay) (h2, s1) (2.00Mbit 2 delay)
(2.00Mbit 2 delay) (h3, s2) (2.00Mbit 2 delay) (2.00Mbit 2 delay) (h4, s2) (8.00Mbit 2 delay 3.00000% loss) (8.00
Mbit 2 delay 3.00000% loss) (s1, s2)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0
*** Starting 2 switches
s1 s2 ... (2.00Mbit 2 delay) (2.00Mbit 2 delay) (8.00Mbit 2 delay 3.00000% loss) (2.00Mbit 2 delay) (2.00Mbit 2 de
lay) (8.00Mbit 2 delay 3.00000% loss)
*** Waiting for switches to connect
s1 s2

Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85.3 KByte (default)

[ 3] local 192.168.1.1 port 58956 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.5 sec   2.50 MBytes  1.99 Mbits/sec

*** Starting CLI:
mininet>
```

## Scheme: cubic with 1% loss

```
mininet@mininet-vm: ~/Downloads
mininet@mininet-vm:~/Downloads$ sudo python3 topology_part2.py --config b --loss 1 --cc cubic
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2
*** Adding links:
(2,00Mbit 2 delay) (2,00Mbit 2 delay) (h1, s1) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h2, s1) (2,00Mbit 2 delay)
(2,00Mbit 2 delay) (h3, s2) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h4, s2) (8,00Mbit 2 delay 1,00000% loss) (8,00
Mbit 2 delay 1,00000% loss) (s1, s2)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0
*** Starting 2 switches
s1 s2 ... (2,00Mbit 2 delay) (2,00Mbit 2 delay) (8,00Mbit 2 delay 1,00000% loss) (2,00Mbit 2 delay) (2,00Mbit 2 de
lay) (8,00Mbit 2 delay 1,00000% loss)
*** Waiting for switches to connect
s1 s2

-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85,3 KByte (default)
-----

[ 3] local 192.168.1.1 port 58976 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0,0-10,7 sec   2,88 MBytes  2,26 Mbits/sec

*** Starting CLI:
mininet>
```

## Scheme: cubic with 3% loss

```
mininet@mininet-vm: ~/Downloads
mininet@mininet-vm:~/Downloads$ sudo python3 topology_part2.py --config b --loss 3 --cc cubic
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2
*** Adding links:
(2,00Mbit 2 delay) (2,00Mbit 2 delay) (h1, s1) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h2, s1) (2,00Mbit 2 delay)
(2,00Mbit 2 delay) (h3, s2) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h4, s2) (8,00Mbit 2 delay 3,00000% loss) (8,00
Mbit 2 delay 3,00000% loss) (s1, s2)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0
*** Starting 2 switches
s1 s2 ... (2,00Mbit 2 delay) (2,00Mbit 2 delay) (8,00Mbit 2 delay 3,00000% loss) (2,00Mbit 2 delay) (2,00Mbit 2 de
lay) (8,00Mbit 2 delay 3,00000% loss)
*** Waiting for switches to connect
s1 s2

-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 85,3 KByte (default)
-----

[ 3] local 192.168.1.1 port 58988 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0,0-10,5 sec   2,62 MBytes  2,10 Mbits/sec

*** Starting CLI:
mininet>
```

## Scheme: bbr with 1% loss

```
mininet@mininet-vm: ~/Downloads
mininet@mininet-vm:~/Downloads$ sudo python3 topology_part2.py --config b --loss 1 --cc bbr
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2
*** Adding links:
(2,00Mbit 2 delay) (2,00Mbit 2 delay) (h1, s1) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h2, s1) (2,00Mbit 2 delay)
(2,00Mbit 2 delay) (h3, s2) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h4, s2) (8,00Mbit 2 delay 1,00000% loss) (8,00
Mbit 2 delay 1,00000% loss) (s1, s2)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0
*** Starting 2 switches
s1 s2 ... (2,00Mbit 2 delay) (2,00Mbit 2 delay) (8,00Mbit 2 delay 1,00000% loss) (2,00Mbit 2 delay) (2,00Mbit 2 de
lay) (8,00Mbit 2 delay 1,00000% loss)
*** Waiting for switches to connect
s1 s2

-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 128 KByte (default)
-----

[ 3] local 192.168.1.1 port 58996 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0,0-10,4 sec   2,50 MBytes  2,02 Mbits/sec

*** Starting CLI:
mininet>
```

## Scheme: bbr with 3% loss

```
mininet@mininet-vm: ~/Downloads
mininet@mininet-vm:~/Downloads$ sudo python3 topology_part2.py --config b --loss 3 --cc bbr
*** Creating network
*** Adding controller
*** Adding hosts:
h1 h2 h3 h4
*** Adding switches:
s1 s2
*** Adding links:
(2,00Mbit 2 delay) (2,00Mbit 2 delay) (h1, s1) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h2, s1) (2,00Mbit 2 delay)
(2,00Mbit 2 delay) (h3, s2) (2,00Mbit 2 delay) (2,00Mbit 2 delay) (h4, s2) (8,00Mbit 2 delay 3,00000% loss) (8,00
Mbit 2 delay 3,00000% loss) (s1, s2)
*** Configuring hosts
h1 h2 h3 h4
*** Starting controller
c0
*** Starting 2 switches
s1 s2 ... (2,00Mbit 2 delay) (2,00Mbit 2 delay) (8,00Mbit 2 delay 3,00000% loss) (2,00Mbit 2 delay) (2,00Mbit 2 de
lay) (8,00Mbit 2 delay 3,00000% loss)
*** Waiting for switches to connect
s1 s2
-----
Client connecting to 192.168.1.4, TCP port 5001
TCP window size: 128 KByte (default)
-----
[ 3] local 192.168.1.1 port 59004 connected with 192.168.1.4 port 5001
[ ID] Interval      Transfer    Bandwidth
[ 3] 0.0-10.3 sec  2.50 MBytes  2.03 Mbits/sec

*** Starting CLI:
mininet>
```

Summary of the results: (h1→h4)

Scheme	Throughput	Transferred Data
<b>RENO (with no loss)</b>	8.03 Mbits/sec	9.88 MB
<b>RENO (with 1% loss)</b>	3.96 Mbits/sec	5.25 MB
<b>RENO (with 3% loss)</b>	2 Mbits/sec	2.50 MB
<b>VEGAS (with no loss)</b>	2.13 Mbits/sec	2.62 MB
<b>VEGAS (with 1% loss)</b>	2.14 Mbits/sec	2.62 MB
<b>VEGAS (with 3% loss)</b>	1.99 Mbits/sec	2.50 MB
<b>CUBIC (with no loss)</b>	3.39 Mbits/sec	4.38 MB
<b>CUBIC (with 1% loss)</b>	2.26 Mbits/sec	2.88 MB
<b>CUBIC (with 3% loss)</b>	2.10 Mbits/sec	2.62 MB
<b>BBR (with no loss)</b>	2.04 Mbits/sec	2.50 MB
<b>BBR (with 1% loss)</b>	2.02 Mbits/sec	2.50 MB
<b>BBR (with 3% loss)</b>	2.03 Mbits/sec	2.50 MB

For reno, we can see that due to 1% and 3% loss on the link, the throughput has decreased significantly.

For vegas, we see that due to 3% loss on the link only, throughput has decreased slightly. For 1% loss on the link, throughput is similar.

For cubic, we see that due to 1% and 3% loss on the link, throughput has decreased.

For bbr, we see that 1% and 3% loss on the link plays no significant role.

The reason for the throughputs on the CCAs are similar to what have mentioned for question b.

